

Behavioral Cloning

The goals / steps of this project are the following:

- Use the simulator to collect data of good driving behavior
- Build a convolution neural network in Keras that predicts steering angles from images
- Train and validate the model with a training and validation set
- Test that the model successfully drives around track one without leaving the road
- Summarize the results with a written report

Figure 1: Short example of car driving autonomously on the jungle course.

Rubric Points

Here I will consider the **rubric points** individually and describe how I addressed each point in my implementation.

Files Submitted & Code Quality

1. Submission includes all required files and can be used to run the simulator in autonomous mode My project includes the following files.

- `model.py` containing the script to create and train the model
- `data.py` containing the helper methods to process data, and our generator function
- `plots.py` script to generate figures presented in this writeup
- `drive.py` for driving the car in autonomous mode
- `model.json` containing a summary of the trained convolution neural network
- `model.h5` containing the weights of the trained convolution neural network
- `final_simple_course.mov` a video of our autonomous vehicle driving the simple course
- `final_jungle_course.mov` a video of our autonomous vehicle driving the jungle course (though with some errors/crashes)
- `jungle_overview.mov` a top-down video of our vehicle navigating a portion of the jungle course
- `writeup.md` and `writeup.pdf`: this document, summarizing the results

2. Submission includes functional code Using the Udacity provided simulator and my `drive.py` file, the car can be driven autonomously around the track by executing the following command. First ensure the proper environment is activated.

```
python drive.py model
```

We trained our dataset on the lowest resolution setting, and so driving on those settings will best reproduce our results. This choice was made to make it easier and quicker to iterate and test our results on our hardware. It has the side effect of making our task a little bit easier, as the courses do not have shadows.

The car should perform flawlessly on the simple track. It can perform well on most of the jungle track but there are a few locations where it fails or gets stuck that we weren't able to resolve before submission.

3. Submission code is usable and readable The `model.py` file contains the code for training and saving the convolution neural network. The file shows the pipeline I used for training and validating the model, and it contains comments to explain how the code works.

Model Architecture and Training Strategy

1. An appropriate model architecture has been employed My model is a straightforward implementation of the convolutional network architecture described in [Bojarski et al. \(2016\)](#), “End-to-End Deep Learning for Self-Driving Cars”. The model consists of five convolutional layers with filter sizes 5x5 and 3x3, followed by three fully-connected layers. The model description and implementation can be found in `model.py`.

Before being fed through the network, images are preprocessed. Preprocessing steps can be found in `data.py` and consist of the following. Images are

1. Converted to a [YUV color space](#),
2. Cropped at top and bottom to remove sky and car artifacts from the images,
3. Downsampled to a resolution that the Nvidia architecture expects (3 color channels at 66x200 pixels), and

The model includes [rectified linear unit (ReLU)]([https://en.wikipedia.org/wiki/Rectifier_\(neural_networks\)](https://en.wikipedia.org/wiki/Rectifier_(neural_networks))) activation function layers to introduce nonlinearity, and the Keras model includes a Lambda layer to normalize pixel brightness values so that they fall in a range of -0.5 to 0.5 instead of 0 to 255.

2. Attempts to reduce overfitting in the model In an attempt to reduce overfitting and encourage the model to generalize, we include three regularizing [Dropout layers](#) after the fully-connected layers. We use a drop rate of 25%. This seemed to have a beneficial though not substantial effect on our network’s performance.

The training data was also collected in an attempt to increase the generalization ability of the trained model. We generated training data driving the simple course forwards and backwards, and collected training data for the jungle course.

We randomly vertically flipped (left turn turns into a right turn) a fraction of the training data during training, flipping also the sign of the ground truth steering measurement.

We used left and right cameras to augment the ground truth data from the center camera, adjusting the steering value by a small amount for the non-centered cameras.

We tested the model by running it (described above) through the driving simulator and ensuring that the vehicle stayed on the test track.

3. Model parameter tuning The model used an [Adam optimizer](#), so the learning rate was not tuned manually.

We tuned the number of epochs, the steering adjustment value applied to non-center camera images, and a factor adjusting how aggressively we balanced the dataset by sampling different steering angles to equalize the steering angle prevalence in the training set.

4. Appropriate training data Training data was collected from the simple and jungle tracks. We drove the car carefully through the courses several times, drove the simple course backwards (to help unbalance the steering angles, as the simple track is biased to left turns), and reinforced difficult regions of the tracks by adding additional training data using proper driving technique.

For details about how I created the training data, see the next section.

Model Architecture and Training Strategy

1. Solution Design Approach The convolutional neural network that we used had to be good for image recognition. A good starting place would be to use any of the many architectures published that have performed well on image-based datasets (see [these suggestions](#)).

However, since this problem specifically had been addressed in the literature before, we used [Bojarski et al. \(2016\)](#)'s network architecture as a starting point. Their architecture was designed specifically for an end-to-end self-driving car. The network consists of several 3x3 and 5x5 convolutional filters, and several fully connected layers.

Our collected data was split into training and validation sets, with 20% of the data devoted to a validation set. By examining the mean squared error value for the training and validation sets over multiple epochs, we could assess how well the model was fitting the data, and if the model was overfitting the data. See the figure below to see accuracy curves over epochs.




Figure 2: Accuracy of training and validation during our training process over all epochs.

If, after many epochs, the training loss continues to improve but the validation loss gets worse, we know that we are beginning to overfit the data, and need to stop training earlier to avoid overfitting.

Since the training data--especially for the simple course--contained most steering values near zero, and more negative (left-turning) turn values than right turn values, we decided to investigate balancing our dataset.

If the training dataset has an overwhelming number of samples with a particular ground truth value, that value may be overpredicted when using the model for prediction. See the following histogram of steering values in our training data below.




Figure 3: Unbalanced Dataset: Histogram showing the frequency of steering angles in the training dataset.

To alleviate this issue, we attempt to more equally represent different steering values in our training data set. To accomplish this, we sample from our training data, discarding some of the samples with over-represented steering values. This reduces the overall size of our training set compared to the collected data, but more equally represents reasonable steering values. This helps ensure, when predicting with our model, non-zero steering values are not suppressed due to their overabundance in the training data.

After equalizing our training data based on steering value, the same graph of steering angle incidence in our training dataset can be seen in the figure below.




Figure 4: Balanced Dataset: Histogram showing the frequency of steering angles after subsampling our training dataset to equalize steering angle values across the entire range.

2. Final Model Architecture The final model architecture was based on the Nvidia convolutional neural network described above, with added dropout layers for regularization. The full network can be seen in the table and figure below.

Layer Type	Output Shape	Number of Parameters
lambda_1 (Lambda)	(66, 200, 3)	0
convolution2d_1 (Convolution2D)	(31, 98, 24)	1824
convolution2d_2 (Convolution2D)	(14, 47, 36)	21636
convolution2d_3 (Convolution2D)	(5, 22, 48)	43248
convolution2d_4 (Convolution2D)	(3, 20, 64)	27712
convolution2d_5 (Convolution2D)	(1, 18, 64)	36928
flatten_1 (Flatten)	(1152)	0
dropout_1 (Dropout)	(1152)	0
dense_1 (Dense)	(100)	115300
dropout_2 (Dropout)	(100)	0
dense_2 (Dense)	(50)	5050
dropout_3 (Dropout)	(50)	0
dense_3 (Dense)	(10)	510
dense_4 (Dense)	(1)	11




Figure 5: Network architecture. We started with the architecture presented in [Bojarski et al. \(2016\)](#)

3. Creation of the Training Set & Training Process We captured data by driving several whole laps in the simple and jungle courses. We also drove the simple course backwards to help unbias the left-leaning nature of the simple course in our training data. After collecting this data, we reinforced certain sections of the track that were difficult for the model to learn by re-driving these sections repeatedly.

Some examples of views in the simple and jungle courses, along with the predicted steering angles associated with these frames, can be seen below.



Figure 6: Example images from our dataset, annotated with the ground truth steering angle.



Figure 7: Example images from our dataset, annotated with the ground truth steering angle.

In the end, our model drove the simple course flawlessly. It was able to drive most of the jungle course but there were a few portions of the track on which we weren't able to successfully train the model.